

# Sentiment Classification DNNs

August 6, 2020

Sentiment with Deep Neural Networks

## 0.1 Outline

By completing this, you will:

- Understand how you can build/design a model using layers
- Train a model using a training loop
- Use a binary cross-entropy loss function
- Compute the accuracy of your model
- Predict using your own input

1

May you've implemented Logistic regression and Naive Bayes for sentiment analysis. However if you were to give your old models an example like:

This movie was almost good. Your model would have predicted a positive sentiment for that review.

However, that sentence has a negative sentiment and indicates that the movie was not good. To solve those kinds of misclassifications, you will write a program that uses deep neural networks to identify sentiment in text. By completing this assignment, you will:

- Understand how you can build/design a model using layers
- Train a model using a training loop
- Use a binary cross-entropy loss function
- Compute the accuracy of your model
- Predict using your own input
- Indeed most of the deep nets you will be implementing will have a similar structure. The only thing

that changes is the model architecture, the inputs, and the outputs. Before starting the assignment, we will introduce you to the Google library trax that we use for building and training models.

Now we will show you how to compute the gradient of a certain function  $f$  by just using `.grad(f)`.

- Trax source code can be found on Github: [Trax](#)
- The Trax code also uses the JAX library: [JAX](#)

## Part 1: Import libraries and try out Trax

In [ ]: `<a name="1"></a>`

*# Part 1: Import libraries and try out Trax*

- Let's import libraries and look at an example of using the Trax library.

In [184]: `import os`

`import random as rnd`

*# import relevant libraries* `import`  
`trax`

*# set random seeds to make this notebook easier to replicate*  
`trax.supervised.trainer_lib.init_random_number_generators(31)`

*# import trax.fastmath.numpy* `import trax.fastmath.numpy as np` *# to make compatible with JAX*

*# import trax.layers* `from trax`  
`import layers as tl`

*# import Layer from the utils.py file* `from utils import Layer,`  
`load_tweets, process_tweet` *# from utils import*

2

### 1 Create an array using `trax.fastmath.numpy`

```
a = np.array(5.0)
```

### 2 View the returned array

```
display(a)
```

```
print(type(a))
```

 Notice that `trax.fastmath.numpy` returns a `DeviceArray` from the `jax` library.

```
In [186]: # Define a function that will use the trax.fastmath.numpy array
```

```
def f(x):
```

```
    #  $f = x^2$  return
```

```
    (x**2)
```

```
In [187]: # Call the function
```

```
print(f'f(a) for a={a} is {f(a)}')
```

```
f(a) for a=5.0 is 25.0
```

The gradient (derivative) of function  $f$  with respect to its input  $x$  is the derivative of  $x^2$ . - The derivative of  $x^2$  is  $2x$ . - When  $x$  is 5, then  $2x = 10$ .

You can calculate the gradient of a function by using `trax.fastmath.grad(fun=)` and passing in the name of the function. - In this case the function you want to take the gradient of is  $f$ . - The object returned (saved in `grad_f` in this example) is a function that can calculate the gradient of  $f$  for a given `trax.fastmath.numpy` array.

```
In [188]: # Directly use trax.fastmath.grad to calculate the gradient (derivative) of the function
```

```
grad_f = trax.fastmath.grad(fun=f) #  $df/dx$  - Gradient of function  $f(x)$  with respect to  $x$ 
```

```
# View the type of the returned object (it 's a function) type(grad_f)
```

```
Out[188]: function
```

```
In [189]: # Call the newly created function and pass in a value for x (the DeviceArray stored in 'a')
```

```
grad_calculation = grad_f(a)
```

```
# View the result of calling the grad_f function|
```

```
display(grad_calculation)
```

```
DeviceArray(10., dtype=float32)
```

The function returned by `trax.fastmath.grad` takes in `x=5` and calculates the gradient of `f`, which is `2*x`, which is 10. The value is also stored as a `DeviceArray` from the `jax` library.

## Part 2: Importing the data

# Part 2: Importing the data ## 2.1 Loading in the data Import the data - Details of `process_tweet` function are available in `utils.py` file

In [190]: ## DO NOT EDIT THIS CELL

```
# Import functions from the utils.py file
```

```
import numpy as np
```

```
# Load positive and negative tweets all_positive_tweets,
all_negative_tweets = load_tweets()
```

```
# View the total number of positive and negative tweets. print(f"The number of positive
tweets: {len(all_positive_tweets)}") print(f"The number of negative tweets:
{len(all_negative_tweets)}")
```

```
# Split positive set into validation and training val_pos = all_positive_tweets[4000:] # generating validation set
for positive tweets train_pos = all_positive_tweets[:4000] # generating training set for positive tweets
```

```
# Split negative set into validation and training val_neg = all_negative_tweets[4000:] # generating validation
set for negative tweets train_neg = all_negative_tweets[:4000] # generating training set for negative tweets
```

```
# Combine training data into one set train_x =
train_pos + train_neg
```

```
# Combine validation data into one set val_x =
val_pos + val_neg
```

```
# Set the labels for the training set (1 for positive, 0 for negative) train_y =
np.append(np.ones(len(train_pos)), np.zeros(len(train_neg)))
```

```
# Set the labels for the validation set (1 for positive, 0 for negative) val_y =
np.append(np.ones(len(val_pos)), np.zeros(len(val_neg)))

print(f"length of train_x {len(train_x)}") print(f"length
of val_x {len(val_x)}")
```

The number of positive tweets: 5000 The  
number of negative tweets: 5000 length of  
train\_x 8000 length of val\_x 2000

4

Now import a function that processes tweets (we’ve provided this in the utils.py file). - ‘process\_tweets’ removes unwanted characters e.g. hashtag, hyperlinks, stock tickers from tweet. - It also returns a list of words (it tokenizes the original string).

```
In [191]: # Import a function that processes the tweets
          # from utils import process_tweet

          # Try out function that processes tweets print("original
          tweet at training position 0") print(train_pos[0])

          print("Tweet at training position 0 after processing:")
          process_tweet(train_pos[0])
```

original tweet at training position 0 #FollowFriday @France\_Inte @PKuchly57 @Milipol\_Paris for being top engaged mem  
community this Tweet at training position 0 after processing:

```
Out[191]: ['followfriday', 'top', 'engag', 'member', 'commun', 'week', ':']
```

Notice that the function process\_tweet keeps key words, removes the hash # symbol, and ignores usernames (words that begin with ‘@’). It also returns a list of the words.

**## 2.2 Building the vocabulary Now build the vocabulary.** - Map each word in each tweet to an integer (an “index”). - The following code does this for you, but please read it and understand what it’s doing. - Note that you will build the vocabulary based on the training data. - To do so, you will assign an index to everyword by iterating over your training set.

The vocabulary will also include some special tokens - \_\_PAD\_\_: padding - </e>: end of line - \_\_UNK\_\_: a token representing any word that is not in the vocabulary.

```

In [ ]: # Build the vocabulary

# Unit Test Note - There is no test set here only train/val

# Include special tokens # started with pad, end of line and unk
tokens Vocab = {'__PAD__': 0, '__</e>__': 1, '__UNK__': 2}

# Note that we build vocab using training data for tweet in
train_x:
    processed_tweet = process_tweet(tweet) for
    word in processed_tweet:
        if word not in Vocab:
            Vocab[word] = len(Vocab)

print("Total words in vocab are",len(Vocab))
display(Vocab)

```

The dictionary Vocab will look like this:

5

```

{'__PAD__': 0,
 '__</e>__': 1, '__UNK__': 2, 'followfriday': 3, 'top': 4, 'engag': 5,

```

- Each unique word has a unique integer associated with it.
- The total number of words in Vocab: 9088

**## 2.3 Converting a tweet to a tensor** Write a function that will convert each tweet to a tensor (a list of unique integer IDs represent- ing the processed tweet). - Note, the returned data type will be a **regular Python list()** - You won't use TensorFlow in this function - You also won't use a numpy array - You also won't use trax.fastmath.numpy array - For words in the tweet that are not in the vocabulary, set them to the unique ID for the token \_\_UNK\_\_.

**Example** Input a tweet:

'@happypuppy, is Maria happy?'

The tweet\_to\_tensor will first conver the tweet into a list of tokens (including only relevant words)

['maria', 'happi']

Then it will convert each word into its unique integer

[2, 56]

- Notice that the word “maria” is not in the vocabulary, so it is assigned the unique integer associated with the \_\_UNK\_\_ token, because it is considered “unknown.”

### Exercise 01 **Instructions:** Write a program `tweet_to_tensor` that takes in a tweet and converts it to an array of numbers. You can use the Vocab dictionary you just found to help create the tensor.

- Use the `vocab_dict` parameter and not a global variable.
- Do not hard code the integer value for the \_\_UNK\_\_ token.

Hints Map each word in tweet to corresponding token in ‘Vocab’ Use Python’s `Dictionary.get(key,value)` so that the function returns a default value if the key is not found in the dictionary.

In [193]: *# UNQ\_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)*

```
# GRADED FUNCTION: tweet_to_tensor def tweet_to_tensor(tweet, vocab_dict,
unk_token='__UNK__', verbose=False):
```

```
    '''
```

6

*Input: tweet - A string containing a tweet*

*vocab\_dict - The words dictionary unk\_token - The special string for*

*unknown tokens verbose - Print info during runtime Output:*

*tensor\_l - A python list with*

```
    '''
```

```
    ### START CODE HERE (Replace instances of 'None' with your code) ###
    # Process the tweet into a list of words # where only important words are kept (stop words
    removed) word_l = process_tweet(tweet)
```

```
    if verbose:
```

```
        print("List of words from the processed tweet:") print(word_l)
```

```
    # Initialize the list that will contain the unique integer IDs of each word tensor_l = []
```

```
    # Get the unique integer ID of the __UNK__ token unk_ID = 2
```

```

if verbose:
    print(f"The unique integer ID for the unk_token is {unk_ID}")

#for each word in the list: for word
in word_l:

    # Get the unique integer ID. # If the word doesn 't exist in the vocab
    dictionary, # use the unique ID for __UNK__ instead. word_ID =
    vocab_dict.get(word, unk_ID) ### END CODE HERE ###

    # Append the unique integer ID to the tensor list.
    tensor_l.append(word_ID)

return tensor_l

```

```

In [194]: print("Actual tweet is\n", val_pos[0])
          print("\nTensor of tweet:\n", tweet_to_tensor(val_pos[0], vocab_dict=Vocab))

```

Actual tweet is

Bro:U wan cut hair anot,ur hair long Liao bo Me:since ord liao,take it easy lor treat as  
save \$ leave it longer :)

7

Bro:LOL SibeI xialan

Tensor of tweet:

[1065, 136, 479, 2351, 745, 8148, 1123, 745, 53, 2, 2672, 791, 2, 2, 349, 601, 2, 3489, 1017, 597, 4559]

### Expected output

Actual tweet is

Bro:U wan cut hair anot,ur hair long Liao bo Me:since ord liao,take it easy lor treat as  
save \$ leave it longer :) Bro:LOL SibeI xialan

Tensor of tweet:

[1065, 136, 479, 2351, 745, 8148, 1123, 745, 53, 2, 2672, 791, 2, 2, 349, 601, 2, 3489, 1017, 597, 4559]

In [195]: # test tweet\_to\_tensor



```

def test_tweet_to_tensor():
    test_cases = [
        {
            "name": "simple_test_check", "input": [val_pos[1], Vocab], "expected": [444, 2, 304, 567, 56,
            9], "error": "The function gives bad output for val_pos[1]. Test failed" }, {
            "name": "datatype_check", "input": [val_pos[1], Vocab], "expected": type([]),
            "error": "Datatype mismatch. Need only list not np.array" }, {
            "name": "without_unk_check", "input": [val_pos[1], Vocab], "expected": 6, "error": "Unk word check not d
            check if you included mapping for unknown } ]
    count = 0
    for test_case in test_cases:

        try:
            if test_case['name'] == "simple_test_check":
                assert test_case["expected"] == tweet_to_tensor(*test_case['input'])
                count += 1
            if test_case['name'] == "datatype_check":

                8
                assert isinstance(tweet_to_tensor(*test_case['input']), test_case["expected"])
                count += 1
            if test_case['name'] == "without_unk_check":
                assert None not in tweet_to_tensor(*test_case['input'])
                count += 1

        except:

    print(test_case['error'])
    if count == 3:

        print("\033[92m All tests passed")
    else:
        print(count, " Tests
        passed out of 3")
    test_tweet_to_tensor()

```

All tests passed

**## 2.4 Creating a batch generator** Most of the time in Natural Language Processing, and AI in general we use batches when training our data sets. - If instead of training with batches of examples, you were to train a model with one example at a time, it would take a very long time to train the model. - You will now build a data generator that takes in the positive/negative tweets and returns a batch of training examples. It returns the model inputs, the targets (positive or negative labels) and the weight for each target (ex: this allows us to can treat some examples as more important to get right than others, but commonly this will all be 1.0).

Once you create the generator, you could include it in a for loop

```
for batch_inputs, batch_targets, batch_example_weights in data_generator:
```

```
...
```

You can also get a single batch like this:

```
batch_inputs, batch_targets, batch_example_weights = next(data_generator)
```

The generator returns the next batch each time it's called. - This generator returns the data in a format (tensors) that you could directly use in your model. - It returns a triple: the inputs, targets, and loss weights: - Inputs is a tensor that contains the batch of tweets we put into the model. - Targets is the corresponding batch of labels that we train to generate. - Loss weights here are just 1s with same shape as targets. Next week, you will use it to mask input padding.

### Exercise 02 Implement data\_generator.

In [196]: # UNQ\_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)

```
# GRADED: Data generator def data_generator(data_pos, data_neg, batch_size, loop, vocab_dict,
shuffle=False):
```

```
    import numpy as np ''' Input: data_pos - Set
```

```
    of positive examples
```

9

```
    data_neg - Set of negative examples batch_size - number of samples per batch. Must be even
    loop - True or False vocab_dict - The words dictionary shuffle - Shuffle the data order
```

```
    Yield: inputs - Subset of positive and negative examples targets - The corresponding labels for
    the subset example_weights - An array specifying the importance of each example
```

```
    ''' ### START GIVEN CODE ###
```

```
        # make sure the batch size is an even number # to allow an equal number of
        positive and negative samples assert batch_size % 2 == 0
```

```
        # Number of positive examples in each batch is half of the batch size # same with number
        of negative examples in each batch n_to_take = batch_size // 2
```

```
        # Use pos_index to walk through the data_pos array # same with
        neg_index and data_neg pos_index = 0 neg_index = 0
```

```
        len_data_pos = len(data_pos)
```

```

len_data_neg = len(data_neg)

# Get an array with the data indexes pos_index_lines
= list(range(len_data_pos)) neg_index_lines =
list(range(len_data_neg))

# shuffle lines if shuffle is set to True if shuffle:
    rnd.shuffle(pos_index_lines)
    rnd.shuffle(neg_index_lines)

stop = False

# Loop indefinitely
while not stop:

    # create a batch with positive and negative examples batch = []

    # First part: Pack n_to_take positive examples

    # Start from pos_index and increment i up to n_to_take

                                1
                                0

    for i in range(n_to_take):

        # If the positive index goes past the positive dataset length, if pos_index >=
        len_data_pos:

            # If loop is set to False, break once we reach the end of the dataset if not loop:
            stop = True;
            break;

        # If user wants to keep re-using the data, reset the index pos_index = 0

        if shuffle:
            # Shuffle the index of the positive sample
            rnd.shuffle(pos_index_lines)

        # get the tweet as pos_index tweet =
        data_pos[pos_index_lines[pos_index]]

```

```

# convert the tweet into tensors of integers representing the processed words tensor =
tweet_to_tensor(tweet, vocab_dict)

# append the tensor to the batch list
batch.append(tensor)

# Increment pos_index by one
pos_index = pos_index + 1

### END GIVEN CODE
###

### START CODE HERE (Replace instances of 'None ' with your code) ###

# Second part: Pack n_to_take negative examples

# Using the same batch list, start from neg_index and increment i up to n_to_take for i in
range(n_to_take):

    # If the negative index goes past the negative dataset length, if
    neg_index >= len_data_neg:

        # If loop is set to False, break once we reach the end of the dataset if not loop:
        stop = True;
        break;

    # If user wants to keep re-using the data, reset the index

        1
        1
    neg_index = 0

    if shuffle:
        rnd.shuffle(neg_index_lines) # Shuffle the index of the
        negative sample

# get the tweet as neg_index tweet =
data_neg[neg_index_lines[neg_index]]

```

```

# convert the tweet into tensors of integers representing the processed words tensor =
tweet_to_tensor(tweet, vocab_dict)

# append the tensor to the batch list
batch.append(tensor)

# Increment neg_index by one
neg_index += 1

### END CODE HERE
###

### START GIVEN CODE ###
    if stop:
        break;

# Update the start index for positive data # so that it 's n_to_take positions after
the current pos_index pos_index += n_to_take

# Update the start index for negative data # so that it 's n_to_take positions after
the current neg_index neg_index += n_to_take

# Get the max tweet length (the length of the longest tweet) # (you will pad all
shorter tweets to have this length) max_len = max([len(q) for q in batch])

# Initialize the input_l, which will # store the padded versions of
the tensors tensor_pad_l = [] # Pad shorter tweets with zeros for
tensor in batch: ### END GIVEN CODE ###

### START CODE HERE (Replace instances of 'None ' with your code) ###
    # Get the number of positions to pad for this tensor so that it will be max_len lo n_pad = max_
len(tensor)

    1
    2

# Generate a list of zeros, with length n_pad pad_l = [0
for i in range(n_pad)] #print(pad_l)

```

```

        # concatenate the tensor and the list of padded zeros tensor_pad = tensor
        + pad_1 # append the padded tensor to the list of padded tensors
        tensor_pad_1.append(tensor_pad)

    # convert the list of padded tensors to a numpy array # and store this as the model inputs inputs
    = np.array(tensor_pad_1) # Generate the list of targets for the positive examples (a list of ones) #
    The length is the number of positive examples in the batch target_pos=[1 for i in
    range(n_to_take)] # Generate the list of targets for the negative examples (a list of zeros) # The
    length i the number of negative examples in the batch target_neg = [0 for i in range(n_to_take)]
    # Concatenate the positive and negative targets target_1 =
    np.concatenate((target_pos,target_neg)) #print(target_1)

    # Convert the target list into a numpy array targets =
    np.array(target_1)

    # Example weights: Treat all examples equally importantly.It should return an np.array example_w
    np.ones_like(targets)

### END CODE HERE
###

### GIVEN CODE ###

    # note we use yield and not return yield inputs,
    targets, example_weights

```

Now you can use your data generator to create a data generator for the training data, and another data generator for the validation data.

We will create a third data generator that does not loop, for testing the final accuracy of the model.

In [197]: # Set the random number generator for the shuffle procedure

```

rnd.seed(30)

# Create the training data generator def
train_generator(batch_size, shuffle = False):
    return data_generator(train_pos, train_neg, batch_size, True, Vocab, shuffle)

# Create the validation data generator

```

3

```
def val_generator(batch_size, shuffle = False):
    return data_generator(val_pos, val_neg, batch_size, True, Vocab, shuffle)

# Create the validation data generator def
test_generator(batch_size, shuffle = False):
    return data_generator(val_pos, val_neg, batch_size, False, Vocab, shuffle)

# Get a batch from the train_generator and inspect. inputs, targets, example_weights =
next(train_generator(4, shuffle=True))

# this will print a list of 4 tensors padded with zeros print(f'Inputs:
{inputs}') print(f'Targets: {targets}') print(f'Example Weights:
{example_weights}')
```

Inputs: [[2005 4451 3201 9 0 0 0 0 0 0]  
[4954 567 2000 1454 5174 3499 141 3499 130 459 9] [3761 109 136 583  
2930 3969 0 0 0 0 0] [ 250 3761 0 0 0 0 0 0 0 0 0]] Targets: [1 1 0 0]  
Example Weights: [1 1 1 1]

In [198]: *# Test the train\_generator*

```
# Create a data generator for training data, # which produces batches of size 4 (for tensors and
their respective targets) tmp_data_gen = train_generator(batch_size = 4)

# Call the data generator to get one batch and its targets tmp_inputs, tmp_targets,
tmp_example_weights = next(tmp_data_gen)

print(f'The inputs shape is {tmp_inputs.shape}') for i,t in
enumerate(tmp_inputs):
    print(f'input tensor: {t}; target {tmp_targets[i]}; example weights {tmp_example_weights[i]')
```

The inputs shape is (4, 14) input tensor: [3 4 5 6 7 8 9 0 0 0 0 0 0 0]; target 1; example weights 1 input tensor: [10 11 12 13  
18 19 20 9 21 22]; target 1; example weights 1 input tensor: [5738 2901 3761 0 0 0 0 0 0 0 0 0 0 0]; target 0; example input  
858 256 3652 5739 307 4458 567 1230 2767 328 1202 3761 0 0]; target 0; example

### Expected output

The inputs shape is (4, 14) input tensor: [3 4 5 6 7 8 9 0 0 0 0 0 0 0]; target 1; example weights 1 input tensor:

[10 11 12 13 14 15 16 17 18 19 20 9 21 22]; target 1; example weights 1

1

4

input tensor: [5738 2901 3761 0 0 0 0 0 0 0 0 0]; target 0; example input tensor: [ 858 256 3652 5739 307 4458 567 123 1202 3761 0 0]; target 0; example

Now that you have your train/val generators, you can just call them and they will return tensors which correspond to your tweets in the first column and their corresponding labels in the second column. Now you can go ahead and start building your neural network.

## Part 3: Defining classes

# Part 3: Defining classes In this part, you will write your own library of layers. It will be very similar to the one used in Trax and also in Keras and PyTorch. Writing your own small framework will help you understand how they all work and use them effectively in the future.

Your framework will be based on the following Layer class from utils.py.

```
class Layer(object):
    """ Base class for layers. """

    # Constructor
    def __init__(self):
        # set weights to None
        self.weights = None

    # The forward propagation should be implemented # by
    # subclasses of this Layer class
    def forward(self, x):
        raise NotImplementedError

    # This function initializes the weights # based on the input signature and
    # random key, # should be implemented by subclasses of this Layer class
    def init_weights_and_state(self, input_signature, random_key):
        pass

    # This initializes and returns the weights, do not override.
    def init(self, input_signature, random_key):
        self.init_weights_and_state(input_signature, random_key)
        return
```



```
self.weights
```

```
# __call__ allows an object of this class # to be
called like it's a function. def __call__(self, x):
    # When this layer object is called, # it calls its
    forward propagation function return self.forward(x)
```

## 3.1 ReLU class You will now implement the ReLU activation function in a class below. The ReLU function looks as follows:

$$\text{ReLU}(x) = \max(0, x)$$

1  
5

### Exercise 03 **Instructions:** Implement the ReLU activation function below. Your function should take in a matrix or vector and it should transform all the negative numbers into 0 while keeping all the positive numbers intact.

Hints Please use `numpy.maximum(A,k)` to find the maximum between each element in A and a scalar k

In [199]: # UNQ\_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)

```
# GRADED FUNCTION:
```

```
Relu class Relu(Layer):
```

```
    """Relu activation function implementation""" def
```

```
    forward(self, x):
```

```
        ''' Input:- x (a numpy array): the input
```

```
        Output:
```

```
        - activation (numpy array): all positive or 0 version of x ''' ### START CODE HERE (Replace
        instances of 'None' with your code) ### #for j in range(len(x)):
```

```
    #for i in range(len(x[j])): # x[j][i]=np.max(x[j][i], 0)
```

```
    x[x<=0]=0 activation = x
```

```
    ### END CODE HERE
```

```
    ###
```

```
    return activation
```

In [200]: # Test your relu function

```
x = np.array([[ -2.0, -1.0, 0.0], [0.0, 1.0, 2.0]], dtype=float) relu_layer = Relu()
```

```
print("Test data is:") print(x) print("Output of Relu is:") print(relu_layer(x))
```

Test data is:  $\begin{bmatrix} -2. \\ -1. \\ 0. \end{bmatrix}$   
 $\begin{bmatrix} 0. & 1. & 2. \end{bmatrix}$  Output of  
Relu is:  $\begin{bmatrix} 0. & 0. & 0. \end{bmatrix}$   
 $\begin{bmatrix} 0. & 1. & 2. \end{bmatrix}$

### Expected Outout

Test data is:  $\begin{bmatrix} -2. \\ -1. \\ 0. \end{bmatrix}$

1  
6

$\begin{bmatrix} 0. & 1. & 2. \end{bmatrix}$  Output of  
Relu is:  $\begin{bmatrix} 0. & 0. & 0. \end{bmatrix}$   
 $\begin{bmatrix} 0. & 1. & 2. \end{bmatrix}$

## ## 3.2 Dense class

### 2.0.1 Exercise

Implement the forward function of the Dense class. - The forward function multiplies the input to the layer (x) by the weight matrix (W)

$$\text{forward}(\mathbf{x}, \mathbf{W}) = \mathbf{xW}$$

- You can use `numpy.dot` to perform the matrix multiplication.

Note that for more efficient code execution, you will use the `trax` version of `math`, which includes a `trax` version of `numpy` and also `random`.

Implement the weight initializer `new_weights` function - Weights are initialized with a random key. - The second parameter is a tuple for the desired shape of the weights (`num_rows`, `num_cols`) - The num of rows for weights should equal the number of columns in `x`, because for forward propagation, you will multiply `x` times weights.

Please use `trax.fastmath.random.normal(key, shape, dtype=tf.float32)` to generate random values for the weight matrix. The key difference between this function and the standard `numpy` randomness is the explicit use of random keys, which need to be passed. While it can look tedious at the first sight to pass the random key everywhere, you will learn in Course 4 why this is very helpful when implementing some

advanced models. - key can be generated by calling `random.get_prng(seed=)` and passing in a number for the seed. - shape is a tuple with the desired shape of the weight matrix. - The number of rows in the weight matrix should equal the number of columns in the variable x. Since x may have 2 dimensions if it represents a single training example (row, col), or three dimensions (batch\_size, row, col), get the last dimension from the tuple that holds the dimensions of x. - The number of columns in the weight matrix is the number of units chosen for that dense layer. Look at the `__init__` function to see which variable stores the number of units. - dtype is the data type of the values in the generated matrix; keep the default of `tf.float32`. In this case, don't explicitly set the dtype (just let it use the default value). Set the standard deviation of the random values to 0.1 - The values generated have a mean of 0 and standard deviation of 1. - Set the default standard deviation stdev to be 0.1 by multiplying the standard deviation to each of the values in the weight matrix.

In [201]: *# use the fastmath module within trax*

```
from trax import fastmath
```

```
# use the numpy module from trax np =  
fastmath.numpy
```

```
# use the fastmath.random module from trax random =  
fastmath.random
```

1  
7

In [202]: *# See how the fastmath.trax.random.normal function works*

```
tmp_key = random.get_prng(seed=1) print("The random seed  
generated by random.get_prng") display(tmp_key)
```

```
print("choose a matrix with 2 rows and 3 columns")  
tmp_shape=(2,3) display(tmp_shape)
```

```
# Generate a weight matrix # Note that you 'll get an error if you try to set dtype to tf.float32, where tf is tensor  
Just avoid setting the dtype and allow it to use the default data type tmp_weight =  
trax.fastmath.random.normal(key=tmp_key, shape=tmp_shape)
```

```
print("Weight matrix generated with a normal distribution with mean 0 and stdev of 1") display(tmp_weight)
```

The random seed generated by `random.get_prng`

```
DeviceArray([0, 1], dtype=uint32)
```

choose a matrix with 2 rows and 3 columns

(2, 3)

Weight matrix generated with a normal distribution with mean 0 and stdev of 1

```
DeviceArray([[ 0.95730704, -0.96992904, 1.0070664 ],  
             [ 0.36619025, 0.17294823, 0.29092228]], dtype=float32)
```

### Exercise 04 Implement  
the Dense class.

In [203]: # *UNQ\_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)*

*# GRADED FUNCTION:  
Dense*

**class** Dense(Layer):

*""" A dense (fully-connected) layer. """*

*# \_\_init\_\_ is implemented for you* **def** *\_\_init\_\_*(self,  
n\_units, init\_stdev=0.1):

1  
8

*# Set the number of units in this layer* **self**.n\_units  
= n\_units **self**.init\_stdev = init\_stdev

*# Please implement 'forward()'* **def**  
**forward**(**self**, x):

### *START CODE HERE (Replace instances of 'None' with your code)* ###

*# Matrix multiply x and the weight matrix* dense =  
np.dot(x, **self**.weights)

```

    ### END CODE HERE
    ###

    return dense

# init_weights def init_weights_and_state(self, input_signature, random_key):

    ### START CODE HERE (Replace instances of 'None' with your code) ###

    # The input_signature has a .shape attribute that gives the shape as a tuple input_shape =
    input_signature.shape shape = (input_shape[-1], self.n_units)

    # Generate the weight matrix from a normal distribution, # and standard
    deviation of 'stdev' w = trax.fastmath.random.normal(random_key,
    shape) w=np.dot(w, self.init_stdev)

    ### END CODE HERE
    ###

    self.weights = w return
    self.weights

```

In [204]: *# Testing your Dense layer*

```

dense_layer = Dense(n_units=10) #sets number of units in dense layer random_key =
random.get_prng(seed=0) # sets random seed z = np.array([[2.0, 7.0, 25.0]]) # input array

dense_layer.init(z, random_key) print("Weights are\n",dense_layer.weights) #Returns randomly generated weights
print("Forward function output is ", dense_layer(z)) # Returns multiplied values of units and weights

```

Weights are

```

[[-0.02837108 0.09368162 -0.10050076 0.14165013 0.10543301 0.09108126 -0.04265672
 0.0986188 -0.05575325 0.00153249] [-0.20785688 0.0554837 0.09142365 0.05744595
 0.07227863 0.01210617
 -0.03237354 0.16234995 0.02450038 -0.13809784]

1
9
[-0.06111237 0.01403724 0.08410042 -0.1094358 -0.10775021 -0.11396459
 -0.05933381 -0.01557652 -0.03832145 -0.11144515]] Forward function output is [[-3.0395496 0.9266802 2.5414743
-2.050473 -1.9769388 -2.582209
 -1.7952735 0.94427425 -0.8980402 -3.7497487 ]]

```

## Expected Outout

Weights are

```
[[-0.02837108 0.09368162 -0.10050076 0.14165013 0.10543301 0.09108126 -0.04265672  
0.0986188 -0.05575325 0.00153249] [-0.20785688 0.0554837 0.09142365 0.05744595  
0.07227863 0.01210617  
-0.03237354 0.16234995 0.02450038 -0.13809784] [-0.06111237 0.01403724 0.08410042  
-0.1094358 -0.10775021 -0.11396459  
-0.05933381 -0.01557652 -0.03832145 -0.11144515]] Foward function output is [[-3.0395496 0.9266802 2.5414743  
-2.050473 -1.9769388 -2.582209  
-1.7952735 0.94427425 -0.8980402 -3.7497487 ]]
```

## 3.3 Model Now you will implement a classifier using neural networks. Here is the model architecture you will be implementing.

For the model implementation, you will use the Trax layers library `tl`. Note that the second character of `tl` is the lowercase of letter L, not the number 1. Trax layers are very similar to the ones you implemented above, but in addition to trainable weights also have a non-trainable state. State is used in layers like batch normalization and for inference, you will learn more about it in course 4.

First, look at the code of the Trax Dense layer and compare to your implementation above. - [tl.Dense](#): Trax Dense layer implementation

One other important layer that you will use a lot is one that allows to execute one layer after another in sequence. - [tl.Serial](#): Combinator that applies layers serially. - You can pass in the layers as arguments to `Serial`, separated by commas. - For example: `tl.Serial(tl.Embeddings(...), tl.Mean(...), tl.Dense(...), tl.LogSoftmax(...))`

Please use the help function to view documentation for each layer.

In [205]: # *View documentation on tl.Dense*

```
#help(tl.Dense)
```

In [206]: # *View documentation on tl.Serial*

```
#help(tl.Serial)
```

- [tl.Embedding](#): Layer constructor function for an embedding layer.

– `tl.Embedding(vocab_size, d_feature)`. – `vocab_size` is the number of unique words in the given vocabulary. – `d_feature` is the number of elements in the word embedding (some choices for a word

embedding size range from 150 to 300, for example). – Recall from the previous course 2, week 4, that the embedding is

In [207]: # *View documentation for tl.Embedding*

```
#help(tl.Embedding)
```

2

0

In [208]: tmp\_embed = tl.Embedding(vocab\_size=3, d\_feature=2)

```
display(tmp_embed)
```

Embedding\_3\_2

- **tl.Mean**: Calculates means across an axis. In this case, please choose axis = 1 to get an average embedding vector (an embedding vector that is an average of all words in the vocabulary).
- For example, if the embedding matrix is 300 elements and vocab size is 10,000 words, taking the mean of the embedding matrix along axis=1 will yield a vector of 300 elements.

In [209]: # *view the documentation for tl.mean*

```
help(tl.Mean)
```

Help on function Mean in module trax.layers.core:

Mean(axis=-1, keepdims=False)

Returns a layer that computes mean values using one tensor axis.

`Mean` uses one tensor axis to form groups of values and replaces each group with the mean value of that group. The resulting values can either remain in their own size 1 axis (`keepdims=True`), or that axis can be removed from the overall tensor (default `keepdims=False`), lowering the rank of the tensor by one.

Args:

axis: Axis along which values are grouped for computing a mean. keepdims: If `True`, keep the resulting size 1 axis as a separate tensor

axis; else, remove that axis.

In [210]: # *Pretend the embedding matrix uses*

```
# 2 elements for embedding the meaning of a word # and has a  
vocabulary size of 3 # So it has shape (2,3) tmp_embed =
```

```
np.array([[1,2,3,],
[4,5,6] ])
```

```
# take the mean along axis 0 print("The mean along axis 0 creates a vector whose length equals the vocabulary size") display(np.mean(tmp_embed,axis=0))
```

```
print("The mean along axis 1 creates a vector whose length equals the number of elements in a word embedding") display(np.mean(tmp_embed,axis=1))
```

2  
1

The mean along axis 0 creates a vector whose length equals the vocabulary size

```
DeviceArray([2.5, 3.5, 4.5], dtype=float32)
```

The mean along axis 1 creates a vector whose length equals the number of elements in a word embedding

```
DeviceArray([2., 5.], dtype=float32)
```

- [tl.LogSoftmax](#): Implements log softmax function
- Here, you don't need to set any parameters for LogSoftMax().

```
In [211]: help(tl.LogSoftmax)
```

Help on function LogSoftmax in module trax.layers.core:

```
LogSoftmax(axis=-1)
```

Returns a layer that applies log softmax along one tensor axis.

`LogSoftmax` acts on a group of values and normalizes them to look like a set of log probability values. (Probability values must be non-negative, and as a set must sum to 1. A group of log probability values can be seen as the natural logarithm function applied to a set of probability values.)

Args:

axis: Axis along which values are grouped for computing log softmax.



## Online documentation

- [tl.Dense](#)
- [tl.Serial](#)
- [tl.Embedding](#)
- [tl.Mean](#)
- [tl.LogSoftmax](#)

### Exercise 05 Implement the **classifier** function.

In [212]: # *UNQ\_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)*

```
# GRADED FUNCTION: classifier
def classifier(vocab_size=len(Vocab), embedding_dim=256,
              output_dim=2, mode='train'):
```

```
### START CODE HERE (Replace instances of 'None ' with your code) ###
```

```

                2
                2
# create embedding layer
embed_layer = tl.Embedding(
    vocab_size=vocab_size, # Size of the vocabulary
    d_feature=embedding_dim) # Embedding dimension

# Create a mean layer, to create an "average" word embedding
mean_layer = tl.Mean(axis=-1, keepdims=False)

# Create a dense layer, one unit for each output
dense_output_layer = tl.Dense(n_units = output_dim)

# Create the log softmax layer (no parameters needed)
log_softmax_layer = tl.LogSoftmax()

# Use tl.Serial to combine all layers # and create
the classifier # of type
trax.layers.combinators.Serial model = tl.Serial(
```

```

tl.Embedding(vocab_size, embedding_dim), # embedding layer tl.Mean(axis=1), #
mean layer tl.Dense(output_dim), # dense output layer tl.LogSoftmax() # log
softmax layer ) ### END CODE HERE ###

```

```

# return the model of type return
model

```

```

In [213]: tmp_model = classifier()

```

```

In [214]: print(type(tmp_model))
display(tmp_model)

```

```

<class 'trax.layers.combinators.Serial'>

```

```

Serial[
Embedding_9088_256
Mean Dense_2
LogSoftmax ]

```

### Expected Outout

```

<class 'trax.layers.combinators.Serial'> Serial[

```

```

2
3

```

```

Embedding_9088_256
Mean Dense_2
LogSoftmax ]

```

### Part 4: Training

# Part 4: Training To train a model on a task, Trax defines an abstraction [trax.supervised.training.TrainTask](#) which packages the train data, loss and optimizer (among other things) together into an object.

Similarly to evaluate a model, Trax defines an abstraction [trax.supervised.training.EvalTask](#) which packages the eval data and metrics (among other things) into another object.

The final piece tying things together is the [trax.supervised.training.Loop](#) abstraction that is a very simple and flexible way to put everything together and train the model, all the while evaluating it and saving checkpoints. Using Loop will save you a lot of code compared to always writing the training loop by hand, like you did in courses 1 and 2. More importantly, you are less likely to have a bug in that code

that would ruin your training.

In [215]: # *View documentation for trax.supervised.training.TrainTask*

```
help(trax.supervised.training.TrainTask)
```

Help on class TrainTask in module trax.supervised.training:

```
class TrainTask(builtins.object)
```

| TrainTask(labeled\_data, loss\_layer, optimizer, lr\_schedule=None, n\_steps\_per\_checkpoint=100) | A supervised task (labeled data + feedback mechanism) for training. | Methods defined here: | `__init__(self, labeled_data, loss_layer, optimizer, lr_schedule=None, n_steps_per_checkpoint=100)` | Configures a training task. | Args: | labeled\_data: Iterator of batches of data tuples. Each tuple has 1+ data (input value) tensors followed by 1 label (target value) tensor. All tensors are NumPy ndarrays or their JAX counterparts. | loss\_layer: Layer that computes a scalar value (the "loss") by comparing | model output  $\hat{y}=f(x)$  to the target  $y$ . | optimizer: Optimizer object that computes model weight updates from | loss-gradients. | lr\_schedule: Learning rate schedule, a function step  $\rightarrow$  learning\_rate. | n\_steps\_per\_checkpoint: How many steps between checkpoints. | `learning_rate(self, step)` | Return the learning rate for the given step. | `next_batch(self)` | Returns one batch of labeled data: a tuple of input(s) plus label.

2  
4

| ----- | Data descriptors defined here: |  
| `__dict__` | dictionary for instance variables (if defined) | `__weakref__` | list of weak references to the object (if defined) | labeled\_data | loss\_layer | n\_steps\_per\_checkpoint | optimizer |  
sample\_batch

In [216]: # *View documentation for trax.supervised.training.EvalTask*

```
help(trax.supervised.training.EvalTask)
```

Help on class EvalTask in module trax.supervised.training:

```
class EvalTask(builtins.object)
```

| EvalTask(labeled\_data, metrics, metric\_names=None, n\_eval\_batches=1) | Labeled data plus scalar functions for (periodically) measuring a model. | An eval task specifies how ('labeled\_data' + 'metrics')

and with what | precision (`n\_eval\_batches`) to measure a model as it is training. | The variance of each scalar output is reduced by measuring over multiple | (`n\_eval\_batches`) batches and reporting the average from those measurements. | Methods defined here: | `__init__(self, labeled_data, metrics, metric_names=None, n_eval_batches=1)` | Configures an eval task: named metrics run with a given data source. | Args: | `labeled_data`: Iterator of batches of labeled data tuples. Each tuple has | 1+ data tensors (NumPy ndarrays) followed by 1 label (target value) | tensor. | `metrics`: List of layers; each computes a scalar value per batch by | comparing model output :math:`\hat{y}=f(x)` to the target :math:`y`.

2  
5

| `metric_names`: List of names, one for each item in `metrics`, in matching | order, to be used when recording/reporting eval output. If None, | generate default names using layer names from metrics. | `n_eval_batches`: Integer N that specifies how many eval batches to run; | the output is then the average of the outputs from the N batches. | `next_batch(self)` | Returns one batch of labeled data: a tuple of input(s) plus label. | \_\_\_\_\_ | Data descriptors defined here: | `__dict__` | dictionary for instance variables (if defined) | `__weakref__` | list of weak references to the object (if defined) | `labeled_data` | `metric_names` | `metrics` | `n_eval_batches` | `sample_batch`

In [217]: # *View documentation for `trax.supervised.training.Loop`*

`help(trax.supervised.training.Loop)`

Help on class Loop in module trax.supervised.training:

class Loop(builtins.object)

| `Loop(model, task, eval_model=None, eval_task=None, output_dir=None, checkpoint_at=None, eval_at=None)` | Loop that for a given number of steps to train a supervised model. | The typical supervised training process randomly initializes a model and updates its weights via feedback (loss-derived gradients) from a training | task, by looping through batches of labeled data. The loop can also | be configured to run periodic evals and save intermediate checkpoints. | For speed, the implementation takes advantage of JAX's composable function | transformations (specifically, `jit` and `grad`). It creates JIT-compiled | pure functions derived from the core model; schematically:

| - training variant: `jit(grad(pure_function(model+loss)))` | - evals variant: `jit(pure_function(model+evals))` | In training or evals, these variants are called with explicit arguments for all relevant input data, model weights/state, optimizer slots, and random number seeds: | - batch: labeled data | - model weights/state: trainable weights and input-related state (e.g., as used by batch normalization) | - optimizer slots: weights in the optimizer that evolve during the training process | - random number seeds: JAX PRNG keys to enable high-quality, distributed, repeatable generation of pseudo-random numbers | Methods defined here: | `__init__(self, task, eval_model=None, eval_task=None, output_dir=None, checkpoint_at=None, eval_at=None, new_rng=None)` | Configures a training 'Loop', including random initialization. | Args: | model: Trax layer, representing the core model to be trained. Loss functions and evaluation metrics (if any) are considered to be outside the core model, taking core model output and data labels as their two inputs. | task: Task instance, which defines the training data, loss function, and optimizer to be used in this training loop. | eval\_model: Optional Trax layer, representing model used for evaluation, e.g., with dropout turned off. If None, the training model (model) will be used. | eval\_task: EvalTask instance or None. If None, don't do any evals. | output\_dir: Path telling where to save outputs (evals and checkpoints). Can be None if both 'eval\_task' and 'checkpoint\_at' are None. | checkpoint\_at: Function (integer --> boolean) that says, for step n, whether that step should have its checkpoint saved. If None, the default is periodic checkpointing at 'task.n\_steps\_per\_checkpoint'. | eval\_at: Function (integer --> boolean) that says, for training step n, whether that step should run evals. If None, run when checkpointing. | `new_rng(self)` | Returns a new single-use random number generator (JAX PRNGKey). | `run(self, n_steps=1)` | Runs this training loop for n steps. Optionally runs evals and saves checkpoints at specified points.

| Args: | n\_steps: Stop training after completing n steps. | `run_evals(self, weights=None, state=None)` | Runs and records evals for this training session. | Args: | weights: Current weights from model in training. | state: Current state from model in training. | `save_checkpoint(self, weights=None, state=None, slots=None)` | Saves checkpoint to disk for the current training step. | Args: | weights: Weights from model being trained. | state: State (non-weight parameters) from model being trained. | slots: Updatable weights for the optimizer in this training loop. |

----- | Data descriptors defined here: | `__dict__` | dictionary for instance variables (if defined) | `__weakref__` | list of weak references to the object (if defined) | `current_step` | Returns current step number in this training session. | `eval_model` | Returns the model used for evaluation. | `model` | Returns the model that is training.

In [218]: # *View optimizers that you could choose from*

```
help(trax.optimizers)
```

Help on package trax.optimizers in trax:

NAME trax.optimizers - Optimizers for use with Trax layers.

PACKAGE

CONTENTS

adafactor

2

8

adam base

momentum

optimizers\_test

rms\_prop sm3

FUNCTIONS

NS

opt\_configure(\*args, \*\*kwargs)

FILE /opt/conda/lib/python3.7/site-packages/trax/optimizers/\_\_init\_\_.py

Notice some available optimizers include:

adafactor

adam

momentum

rms\_prop

sm3

**## 4.1 Training the model** Now you are going to train your model. Let's define the TrainTask, EvalTask and Loop in preparation to train the model.

In [219]: **from** trax.supervised **import** training

```
batch_size = 16
rnd.seed(271)

train_task = training.TrainTask(
    labeled_data=train_generator(batch_size=batch_size, shuffle=True),
    loss_layer=tl.CrossEntropyLoss(), optimizer=trax.optimizers.Adam(0.01),
    n_steps_per_checkpoint=10, )eval_task = training.EvalTask(
    labeled_data=val_generator(batch_size=batch_size, shuffle=True),
    metrics=[tl.CrossEntropyLoss(), tl.Accuracy()], )model = classifier()
```

This defines a model trained using `tl.CrossEntropyLoss` optimized with the `trax.optimizers.Adam` optimizer, all the while tracking the accuracy using `tl.Accuracy` metric. We also track `tl.CrossEntropyLoss` on the validation set.

2  
9

Now let's make an output directory and train the model.

In [220]: `output_dir = '~/model/'`

```
output_dir_expand = os.path.expanduser(output_dir)
print(output_dir_expand)
```

`/home/jovyan/model/`

### Exercise 06 **Instructions:** Implement `train_model` to train the model (classifier that you wrote earlier) for the given number of training steps (`n_steps`) using `TrainTask`, `EvalTask` and `Loop`.

In [221]: `# UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)`

```
# GRADED FUNCTION: train_model def train_model(classifier, train_task, eval_task,
n_steps, output_dir):
```

```
    ''' Input: classifier - the model you are building
    train_task - Training task eval_task - Evaluation task
    n_steps - the evaluation steps output_dir - folder to
    save your files Output:
```

```
trainer - trax trainer ''' ### START CODE HERE (Replace instances of 'None' with your
```

```

code) ### training_loop = training.Loop(classifier, # The learning model
                                     train_task, # The training task eval_task=eval_task,
                                     output_dir = output_dir) # The output directory

training_loop.run(n_steps = n_steps) ### END CODE
HERE ###

# Return the training_loop, since it has the model. return
training_loop

```

In [222]: training\_loop = train\_model(model, train\_task, eval\_task, 100, output\_dir\_expand)

```

Step 1: train CrossEntropyLoss | 1.12272775 Step 1: eval
CrossEntropyLoss | 0.89718306 Step 1: eval Accuracy |
0.31250000 Step 10: train CrossEntropyLoss | 0.81427336 Step
10: eval CrossEntropyLoss | 0.50847912 Step 10: eval
Accuracy | 0.75000000 Step 20: train CrossEntropyLoss |
0.52660501 Step 20: eval CrossEntropyLoss | 0.34810668

```

3  
0

```

Step 20: eval Accuracy | 0.87500000 Step 30: train
CrossEntropyLoss | 0.32775387 Step 30: eval
CrossEntropyLoss | 0.30329353 Step 30: eval Accuracy |
0.87500000 Step 40: train CrossEntropyLoss | 0.29091102 Step
40: eval CrossEntropyLoss | 0.30724931 Step 40: eval
Accuracy | 0.87500000 Step 50: train CrossEntropyLoss |
0.17655556 Step 50: eval CrossEntropyLoss | 0.13073292 Step
50: eval Accuracy | 1.00000000 Step 60: train
CrossEntropyLoss | 0.16591090 Step 60: eval
CrossEntropyLoss | 0.13529448 Step 60: eval Accuracy |
1.00000000 Step 70: train CrossEntropyLoss | 0.12349443 Step
70: eval CrossEntropyLoss | 0.07977644 Step 70: eval
Accuracy | 1.00000000 Step 80: train CrossEntropyLoss |
0.11218882 Step 80: eval CrossEntropyLoss | 0.02898891 Step
80: eval Accuracy | 1.00000000 Step 90: train
CrossEntropyLoss | 0.07728098 Step 90: eval
CrossEntropyLoss | 0.02353959 Step 90: eval Accuracy |
1.00000000 Step 100: train CrossEntropyLoss | 0.06835373
Step 100: eval CrossEntropyLoss | 0.00595517 Step 100: eval
Accuracy | 1.00000000

```



### Expected output (Approximately)

```
Step 1: train CrossEntropyLoss | 0.88939196 Step 1: eval  
CrossEntropyLoss | 0.68833977 Step 1: eval Accuracy |  
0.50000000 Step 10: train CrossEntropyLoss | 0.61036736 Step  
10: eval CrossEntropyLoss | 0.52182281 Step 10: eval  
Accuracy | 0.68750000 Step 20: train CrossEntropyLoss |  
0.34137666 Step 20: eval CrossEntropyLoss | 0.20654774 Step  
20: eval Accuracy | 1.00000000 Step 30: train  
CrossEntropyLoss | 0.20208922 Step 30: eval  
CrossEntropyLoss | 0.21594886 Step 30: eval Accuracy |  
0.93750000 Step 40: train CrossEntropyLoss | 0.19611198 Step  
40: eval CrossEntropyLoss | 0.17582777 Step 40: eval  
Accuracy | 1.00000000 Step 50: train CrossEntropyLoss |  
0.11203773 Step 50: eval CrossEntropyLoss | 0.07589275 Step  
50: eval Accuracy | 1.00000000 Step 60: train  
CrossEntropyLoss | 0.09375446
```

3

1

```
Step 60: eval CrossEntropyLoss | 0.09290724 Step 60: eval  
Accuracy | 1.00000000 Step 70: train CrossEntropyLoss |  
0.08785903 Step 70: eval CrossEntropyLoss | 0.09610598 Step  
70: eval Accuracy | 1.00000000 Step 80: train  
CrossEntropyLoss | 0.08858261 Step 80: eval  
CrossEntropyLoss | 0.02319432 Step 80: eval Accuracy |  
1.00000000 Step 90: train CrossEntropyLoss | 0.05699894 Step  
90: eval CrossEntropyLoss | 0.01778970 Step 90: eval  
Accuracy | 1.00000000 Step 100: train CrossEntropyLoss |  
0.03663783 Step 100: eval CrossEntropyLoss | 0.00210550  
Step 100: eval Accuracy | 1.00000000
```

**## 4.2 Practice Making a prediction** Now that you have trained a model, you can access it as `training_loop.model` object. We will actually use `training_loop.eval_model` and in the next weeks you will learn why we sometimes use a different model for evaluation, e.g., one without dropout. For now, make predictions with your model.

Use the training data just to see how the prediction process works. - Later, you will use validation data to evaluate your model's performance.

In [223]: *# Create a generator object*

```

tmp_train_generator = train_generator(16)

# get one batch tmp_batch =
next(tmp_train_generator)

# Position 0 has the model inputs (tweets as tensors) # position 1 has the
targets (the actual labels) tmp_inputs, tmp_targets,
tmp_example_weights = tmp_batch

print(f"The batch is a tuple of length {len(tmp_batch)} because position 0 contains the tweets print(f"The shape
tweet tensors is {tmp_inputs.shape} (num of examples, length of tweet print(f"The shape of the labels is
{tmp_targets.shape}, which is the batch size.") print(f"The shape of the example_weights is
{tmp_example_weights.shape}, which is the same as

```

The batch is a tuple of length 3 because position 0 contains the tweets, and position 1 contains the targets. The shape of the tweet tensors is (16, 15) (num of examples, length of tweet tensors). The shape of the labels is (16, 1), which is the batch size. The shape of the example\_weights is (16, 1), which is the same as inputs/targets size.

In [224]: # feed the tweet tensors into the model to get a prediction

```

tmp_pred = training_loop.eval_model(tmp_inputs) print(f"The prediction shape is {tmp_pred.shape},
num of tensor_tweets as rows") print("Column 0 is the probability of a negative sentiment (class 0)")
print("Column 1 is the probability of a positive sentiment (class 1)")

3
2

print() print("View the prediction array")
tmp_pred

```

The prediction shape is (16, 2), num of tensor\_tweets as rows. Column 0 is the probability of a negative sentiment (class 0). Column 1 is the probability of a positive sentiment (class 1).

View the prediction array

```

Out[224]: DeviceArray([[ -4.3611403e+00,  -1.2845993e-02],  [-4.1929121e+00,
-1.5217304e-02],  [-4.8559332e+00,  -7.8125000e-03],  [-3.6934922e+00,  -2.5199771e-02],
[-3.5967889e+00,  -2.7794361e-02],  [-3.5383053e+00,  -2.9493213e-02],  [-5.2808285e+00,
-5.1012039e-03],  [-5.4767599e+00,  -4.1916370e-03],  [-1.2308240e-02,  -4.4036212e+00],
[-5.5770874e-03,  -5.1918826e+00],  [-8.0459118e-03,  -4.8266072e+00],  [-7.2956085e-05,
-9.5227098e+00],  [-1.6406775e-02,  -4.1182537e+00],  [-2.8779507e-03,  -5.8520947e+00],

```

```
[-1.1285543e-02, -4.4898720e+00], [-4.8916340e-03, -5.3226776e+00]], dtype=float32)
```

To turn these probabilities into categories (negative or positive sentiment prediction), for each row: - Compare the probabilities in each column. - If column 1 has a value greater than column 0, classify that as a positive tweet. - Otherwise if column 1 is less than or equal to column 0, classify that example as a negative tweet.

In [225]: # *turn probabilities into category predictions*

```
tmp_is_positive = tmp_pred[:,1] > tmp_pred[:,0] for i, p in
enumerate(tmp_is_positive):
    print(f"Neg log prob {tmp_pred[i,0]:.4f}\tPos log prob {tmp_pred[i,1]:.4f}\t is positive?
```

```
Neg log prob -4.3611 Pos log prob -0.0128 is positive? True actual 1 Neg log prob -4.1929 Pos log prob -0.0152 is
positive? True actual 1 Neg log prob -4.8559 Pos log prob -0.0078 is positive? True actual 1 Neg log prob -3.6935 Pos
log prob -0.0252 is positive? True actual 1 Neg log prob -3.5968 Pos log prob -0.0278 is positive? True actual 1 Neg log
prob -3.5383 Pos log prob -0.0295 is positive? True actual 1 Neg log prob -5.2808 Pos log prob -0.0051 is positive?
True actual 1 Neg log prob -5.4768 Pos log prob -0.0042 is positive? True actual 1 Neg log prob -0.0123 Pos log prob
-4.4036 is positive? False actual 0 Neg log prob -0.0056 Pos log prob -5.1919 is positive? False actual 0 Neg log prob
-0.0080 Pos log prob -4.8266 is positive? False actual 0
```

3

3

```
Neg log prob -0.0001 Pos log prob -9.5227 is positive? False actual 0 Neg log prob -0.0164 Pos log prob -4.1183 is
positive? False actual 0 Neg log prob -0.0029 Pos log prob -5.8521 is positive? False actual 0 Neg log prob -0.0113 Pos
log prob -4.4899 is positive? False actual 0 Neg log prob -0.0049 Pos log prob -5.3227 is positive? False actual 0
```

Notice that since you are making a prediction using a training batch, it's more likely that the model's predictions match the actual targets (labels). - Every prediction that the tweet is positive is also matching the actual target of 1 (positive sentiment). - Similarly, all predictions that the sentiment is not positive matches the actual target of 0 (negative sentiment)

One more useful thing to know is how to compare if the prediction is matching the actual target (label). - The result of calculation is `is_positive` is a boolean. - The target is a type `trax.fastmath.numpy.int32` - If you expect to be doing division, you may prefer to work with decimal numbers with the data type type `trax.fastmath.numpy.int32`

In [226]: # *View the array of booleans*

```
print("Array of booleans")
display(tmp_is_positive)
```

```
# convert boolean to type int32 # True is converted to 1 # False is
converted to 0 tmp_is_positive_int = tmp_is_positive.astype(np.int32)
```

```

# View the array of integers
print("Array of integers")
display(tmp_is_positive_int)

# convert boolean to type float32 tmp_is_positive_float =
tmp_is_positive.astype(np.float32)

# View the array of floats print("Array
of floats")
display(tmp_is_positive_float)

```

Array of booleans

```

DeviceArray([ True,  True,  True,  True,  True,  True,  True,  True,
             False, False, False, False, False, False, False, False], dtype=bool)

```

Array of integers

```

3
4
DeviceArray([1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int32)

```

Array of floats

```

DeviceArray([1., 1., 1., 1., 1., 1., 1., 1., 0., 0., 0., 0., 0., 0.,
            0.], dtype=float32)

```

In [227]: tmp\_pred.shape

Out[227]: (16, 2)

Note that Python usually does type conversion for you when you compare a boolean to an integer - True compared to 1 is True, otherwise any other integer is False. - False compared to 0 is True, otherwise any other integer is False.

```
In [228]: print(f"True == 1: {True == 1}") print(f"True
== 2: {True == 2}") print(f"False == 0: {False == 0}")
print(f"False == 2: {False == 2}")
```

```
True == 1: True True
== 2: False False ==
0: True False == 2:
False
```

However, we recommend that you keep track of the data type of your variables to avoid unexpected outcomes. So it helps to convert the booleans into integers - Compare 1 to 1 rather than comparing True to 1.

Hopefully you are now familiar with what kinds of inputs and outputs the model uses when making a prediction. - This will help you implement a function that estimates the accuracy of the model's predictions.

# Part 5: Evaluation ## 5.1 Computing the accuracy on a batch You will now write a function that evaluates your model on the validation set and returns the accuracy. - preds contains the predictions. - Its dimensions are (batch\_size, output\_dim). output\_dim is two in this case. Column 0 contains the probability that the tweet belongs to class 0 (negative sentiment). Column 1 contains probability that it belongs to class 1 (positive sentiment). - If the probability in column 1 is greater than the probability in column 0, then interpret this as the model's prediction that the example has label 1 (positive sentiment). - Otherwise, if the probabilities are equal or the probability in column 0 is higher, the model's prediction is 0 (negative sentiment). - y contains the actual labels. - y\_weights contains the weights to give to predictions.

### Exercise 07 Implement compute\_accuracy.

3  
5

```
In [229]: # UNQ_C7 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
```

```
# GRADED FUNCTION: compute_accuracy def
```

```
compute_accuracy(preds, y, y_weights):
```

```
""" Input: preds: a tensor of shape (dim_batch, output_dim)
```

```
y: a tensor of shape (dim_batch, output_dim) with the true labels y_weights: a n.ndarray
with the a weight for each example Output:
```

```
accuracy: a float between 0-1 weighted_num_correct (np.float32): Sum of the weighted correct
predictions sum_weights (np.float32): Sum of the weights """ ### START CODE HERE (Replace
```

```
instances of 'None ' with your code) ### # Create an array of booleans, # True if the probability of
positive sentiment is greater than # the probability of negative sentiment # else False is_pos =
```

```

np.array(preds[:, 1]>preds[:, 0])

#convert the array of booleans into an array of np.int32 is_pos_int =
is_pos.astype(np.int32)

#compare the array of predictions (as int32) with the target (labels) of type int32 correct =
np.array(is_pos_int[:,]==y[:,])

# Count the sum of the weights.
sum_weights = np.sum(y_weights)

#convert the array of correct predictions (boolean) into an array of np.float32 correct_float =
correct.astype(np.float32)

# Multiply each prediction with its corresponding weight. weighted_correct_float =
[np.multiply(correct_float[i],y_weights[i]) for i in range(len(co

# Sum up the weighted correct predictions (of type np.float32), to go in the # denominator.
weighted_num_correct = np.sum(weighted_correct_float)

# Divide the number of weighted correct predictions by the sum of the # weights. accuracy
= weighted_num_correct/sum_weights

### END CODE HERE ### return accuracy,
weighted_num_correct, sum_weights

```

In [230]: # test your function

```

tmp_val_generator = val_generator(64)

3
6

# get one batch tmp_batch =
next(tmp_val_generator)

# Position 0 has the model inputs (tweets as tensors) # position 1 has the
targets (the actual labels) tmp_inputs, tmp_targets,
tmp_example_weights = tmp_batch

#feed the tweet tensors into the model to get a prediction tmp_pred =
training_loop.eval_model(tmp_inputs)

```

```
tmp_acc, tmp_num_correct, tmp_num_predictions = compute_accuracy(preds=tmp_pred, y=tmp_targets
```

```
print(f'Model's prediction accuracy on a single training batch is: {100 * tmp_acc}%') print(f'Weighted number of correct predictions {tmp_num_correct}; weighted number of total observations predicted {tmp_num_predictions}')

```

Model's prediction accuracy on a single training batch is: 100.0% Weighted number of correct predictions 64.0; weighted number of total observations predicted 64

### Expected output (Approximately)

Model's prediction accuracy on a single training batch is: 100.0% Weighted number of correct predictions 64.0; weighted number of total observations predicted 64

## 5.2 Testing your model on Validation Data Now you will write test your model's prediction accuracy on validation data. This program will take in a data generator and your model. - The generator allows you to get batches of data. You can use it with a for loop:

for batch in iterator:

```
# do something with that batch
```

batch has dimensions (batch size, 2). - Column 0 corresponds to the tweet as a tensor. - Column 1 corresponds to its target (actual label, positive or negative sentiment). - You can feed the tweet into model and it will return the predictions for the batch.

### Exercise 08 **Instructions:** - Compute the accuracy over all the batches in the validation iterator. - Make use of compute\_accuracy, which you recently implemented, and return the overall accuracy.

In [231]: # UNQ\_C8 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)

```
# GRADED FUNCTION: test_model def
```

```
test_model(generator, model):
```

```
''' Input: generator: an iterator instance that provides batches of inputs and targets
```

```
model: a model instance Output:
```

```
accuracy: float corresponding to the accuracy
```

```
3
```

```
7
```

```
'''
```

```
,
```

```
accuracy = 0.
```

```
total_num_correct = 0
total_num_pred = 0
```

```
### START CODE HERE (Replace instances of 'None' with your code) ###
for batch in generator:
```

```
    # Retrieve the inputs from the batch
    inputs = None
```

```
    # Retrieve the targets (actual labels) from the batch
    targets = None
```

```
    # Retrieve the example weight.
    example_weight = None
```

```
    # Make predictions using the inputs
    pred = None
```

```
    # Calculate accuracy for the batch by comparing its predictions and targets
    batch_accuracy, batch_num_correct, batch_num_pred = None
```

```
    # Update the total number of correct predictions # by adding the number of
    # correct predictions from this batch
    total_num_correct += None
```

```
    # Update the total number of predictions # by adding the number of
    # predictions made for the batch
    total_num_pred += None
```

```
    # Calculate accuracy over all examples
    accuracy = None
```

```
### END CODE HERE
### return accuracy
```

In [232]: # DO NOT EDIT THIS CELL

```
# testing the accuracy of your model: this takes around 20 seconds
model = training_loop.eval_model
accuracy = test_model(test_generator(16), model)
```

```
print(f'The accuracy of your model on the validation set is {accuracy:.4f}', )
```

-----



```

3
8
TypeError Traceback (most recent call last)

<ipython-input-232-5a6afde371c2> in <module>
2 # testing the accuracy of your model: this takes around 20 seconds 3 model =
training_loop.eval_model ----> 4 accuracy = test_model(test_generator(16), model)

56 print(f'The accuracy of your model on the validation set is {accuracy:.4f}', )

<ipython-input-231-cd65dad7d79d> in test_model(generator, model)
3031 # Calculate accuracy for the batch by comparing its predictions and targets ---> 32 batch_accuracy,
batch_num_correct, batch_num_pred = None

3334 # Update the total number of correct predictions

TypeError: cannot unpack non-iterable NoneType object

```

### Expected Output (Approximately)

The accuracy of your model on the validation set is 0.9931

# Part 6: Testing with your own input Finally you will test with your own input. You will see that deepnets are more powerful than the older methods you have used before. Although you go close to 100% accuracy on the first two assignments, the task was way easier.

In [ ]: *# this is used to predict on your own sentence*

```

def predict(sentence):
    inputs = np.array(tweet_to_tensor(sentence, vocab_dict=Vocab))

    # Batch size 1, add dimension for batch, to work with the model inputs =
    inputs[None, :]

    # predict with the model
    preds_probs = model(inputs)

```

```

# Turn probabilities into categories preds = int(preds_probs[0, 1]
> preds_probs[0, 0])

sentiment = "negative" if
preds == 1:
    sentiment = 'positive'

return preds, sentiment

```

3  
9

In [ ]: # try a positive sentence

```

sentence = "It's such a nice day, think i'll be taking Sid to Ramsgate fish and chips for lunch tmp_pred, tmp_senti
predict(sentence) print(f"The sentiment of the sentence \n***\n\"{sentence}\"'\n***\nis {tmp_sentiment}.")

print() # try a negative sentence sentence = "I hated my day, it was the worst, I'm so sad." tmp_pred,
tmp_sentiment = predict(sentence) print(f"The sentiment of the sentence \n***\n\"{sentence}\"'\n***\nis
{tmp_sentiment}.")

```

Notice that the model works well even for complex sentences.

## 2.0.2 On Deep Nets

Deep nets allow you to understand and capture dependencies that you would have not been able to capture with a simple linear regression, or logistic regression. - It also allows you to better use pre-trained embeddings for classification and tends to generalize better.

4  
0